


DroneMOM

Drone Model Output Machine

Taylor Nelms, John Marcao, Eric Micallef


Overview

Create a tool that converts incoming visual and inertial measurement unit data from an Unmanned Aerial Vehicle and use it to generate glTF models of classified objects in the scene.



Jetson Challenges

- Many Low memory warnings
- Classification networks must be lightweight
- Feature matches using SURF instead of SIFT for performance reasons



Classification

- Accomplished
 - Integrated TensorRT with a Universal File Format
 - Using SSD-MobileNet Architecture
- Challenges
 - Converting incoming IMU data into world position
 - Synchronizing data between pipeline stages
- Next Week
 - Convert IMU into a position for point cloud use
 - Find a suitable ROS bag for demoing
 - Benchmark + Stream TensorRT implementation to prove correctness
 - Create own ROS msg that bundles all messages together in a time synced manner



What is TensorRT?

- Think of it like DXR
- Build a tensorRT engine (like CPU side of DXR)
- Run inference on GPU based on this accelerated engine.

TensorRT Optimizations and Performance



Weight & Activation Precision Calibration

Maximizes throughput by quantizing models to INT8 while preserving accuracy



Layer & Tensor Fusion

Optimizes use of GPU memory and bandwidth by fusing nodes in a kernel



Kernel Auto-Tuning

Selects best data layers and algorithms based on target GPU platform



Dynamic Tensor Memory

Minimizes memory footprint and re-uses memory for tensors efficiently



Multi-Stream Execution

Scalable design to process multiple input streams in parallel

Point Cloud Generation


- Accomplished
 - Learned a lot
 - Discovered the difficulty of working with uncalibrated camera data
 - Got feature matching between subsequent frames using SURF feature detection




Point Cloud Generation

- Challenges
 - Using made-up FoV values lead to issues
- Next Week
 - Feature matches and 3d reconstruction on calibrated camera data

FOV: 90





FOV: 120





GLTF Mesh Construction

- Accomplished
 - Batching of mesh construction jobs to reduce delay
 - Downsample and smoothing before reconstruction to remove unimportant data points
 - Research into possible improvements to mesh reconstruction
- Old
 - 20 Frames, 28MB Mesh, 1318 Seconds
- New
 - 20 Frames, 3MB Mesh, 15 Seconds



GLTF Mesh Construction

Combining more frames data gives better models...



...but takes more time and resources to process.

GLTF Mesh Construction

- Challenges

- Almost every GLTF tool is written in JavaScript.
- Point Cloud Library CUDA support is limited. .
- VTK file format not good for GLTF conversion.

- Next Week

- Finalize GLTF export
 - Using cgltf library, allows for writing mesh to memory and then dumping out to a file.
- Implement Mesh Construction on GPU.
 - Still deciding on implementations best for our platform.

End

Questions?